



# Evolving GitOps: Harnessing Kubernetes Resource Model for 5G Core

Alexander North, Ashan Senevirathne & Joel Studler



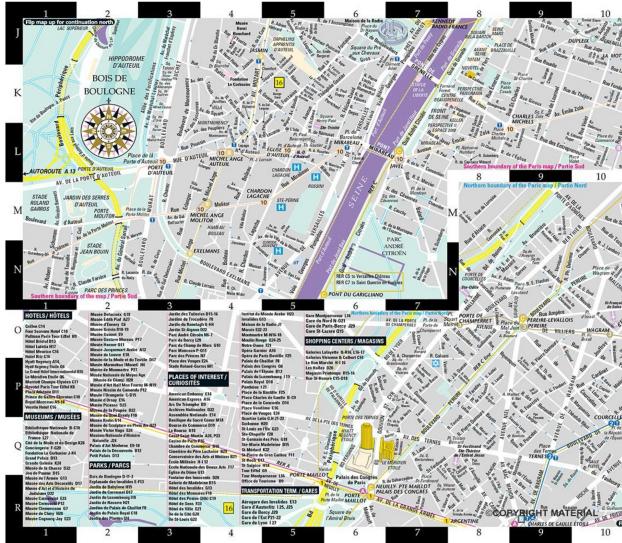


# An Analogy...

## Static paper map

- Fixed
- Static
- Unchanging
- Overwhelming

This is GitOps today



## Apple maps

- Dynamic
- Changes based on external conditions
- More focused
- Simple to navigate

This is GitOps w/ KRM





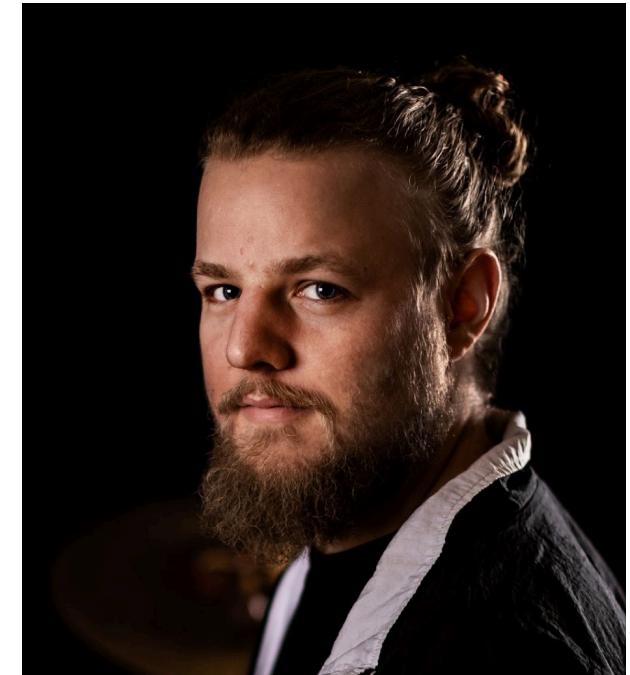
**Alexander North**  
DevOps Engineer

[alexander.north1@swisscom.com](mailto:alexander.north1@swisscom.com)  
Slack: CNCF, Nephio, NetDev



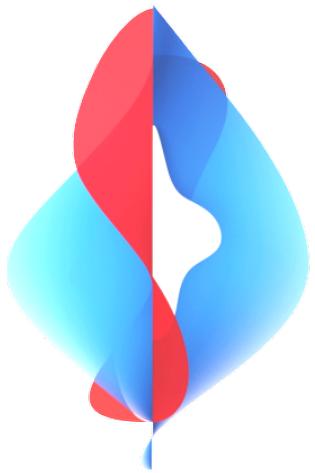
**Ashan Senevirathne**  
Product Owner

[ashan.senevirathne@swisscom.com](mailto:ashan.senevirathne@swisscom.com)  
Slack: CNCF, Nephio, NetDev



**Joel Studler**  
Senior DevOps Engineer

[joel.studler@swisscom.com](mailto:joel.studler@swisscom.com)  
Slack: CNCF, Nephio, NetDev



**swisscom**



# From Telco to TechCo: What Does It Mean for Us?

## Cloud Nativeness

we need to introduce Reconciliation and leverage Kubernetes Patterns

## Simplicity

we need to introduce abstraction in our configuration and break it into smaller manageable parts

## Automation

we need to move from an imperative workflow-driven model to a declarative intent-based model

### 5G - driving our journey from Telco to TechCo

by Swisscom CTIO Mark Düsener at Connect Conference 2022

<https://www.youtube.com/watch?v=hND7TiXJED8>





# What Is a 5G Core?

## Each blue object

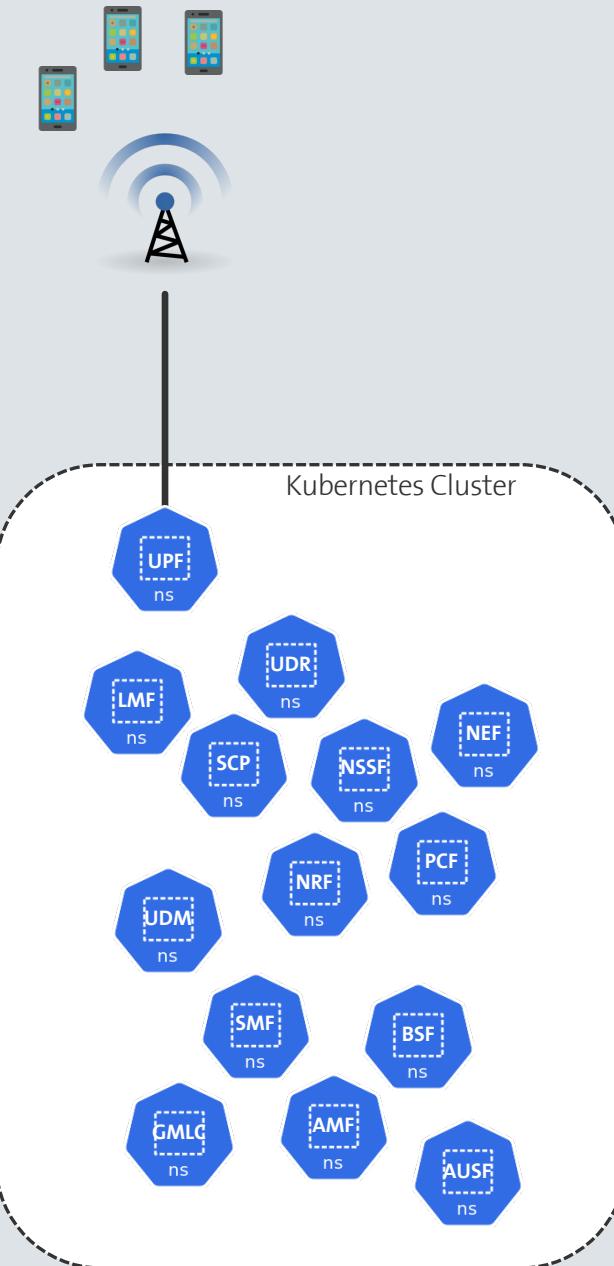
- is a «CNF» aka «Containerized Network Function»
  - e.g. Router (UPF), Authentication Service (AUSF)
- Deployed using Helm

## Configuration is done via

- Helm Values
- Other Configuration Interfaces

## Scale

- A development environment contains ~2000 pods
- A total of 5000 interdependent configuration parameters

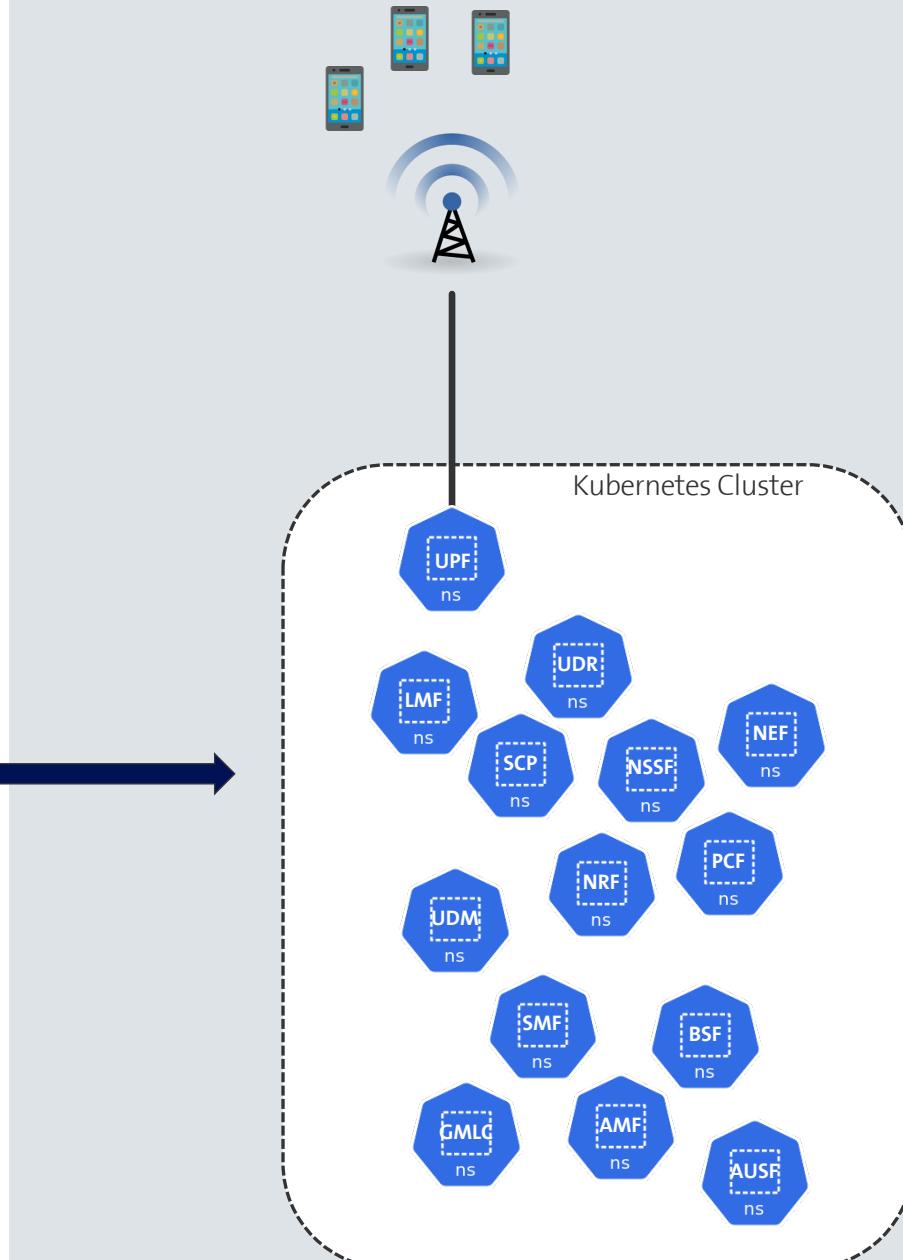




# What Is a 5G Core?

IP addresses  
Subnets  
VLANs  
DNS Records  
Network function variables  
Infrastructure variables  
Network function-Network function mapping  
Secret references  
Certificate references

```
</application>
<application>
  <application-name>slg</application-name>
  <default-load-sharing>true</default-load-sharing>
  <local-host>
    <host-name>afe23</host-name>
    <realm-name>ecp.009.999.mobilenet</realm-name>
  </local-host>
  <realm>
    <realm-name>ecp.009.999.mobilenet</realm-name>
    <peer-list>1</peer-list>
    <peer-list>2</peer-list>
    <realm-load-sharing>true</realm-load-sharing>
  </realm>
</application>
<peers>
  <peer>
    <ipv4v6-address>192.168.244.253</ipv4v6-address>
    <peer-number>1</peer-number>
    <peer-port-number>3868</peer-port-number>
    <is-geographically-redundant>false</is-geographically-redundant>
    <local-host>
      <host-name>afe23</host-name>
      <realm-name>ecp.009.999.mobilenet</realm-name>
    </local-host>
  </peer>
  <peer>
    <ipv4v6-address>192.168.244.213</ipv4v6-address>
    <peer-number>2</peer-number>
    <peer-port-number>3868</peer-port-number>
    <is-geographically-redundant>false</is-geographically-redundant>
    <local-host>
      <host-name>afe23</host-name>
      <realm-name>ecp.009.999.mobilenet</realm-name>
    </local-host>
  </peer>
<local-host>
  <host-name>afe23</host-name>
  <realm-name>ecp.009.999.mobilenet</realm-name>
</local-host>
<sctp-end-point>
  <sctp-end-point-no>1</sctp-end-point-no>
</sctp-end-point>
</local-host>
<diameter>
<dnn-function operation="replace">
  <dnn-redirection-enabled>true</dnn-redirection-enabled>
  <dnn-resolution-extension-enabled>false</dnn-resolution-extension-enabled>
</dnn-function>
<dnn-redirection-profile operation="replace">
  <dnn-redirection-profile-name>DefaultDNN</dnn-redirection-profile-name>
  <dnn-redirection-rule>defaultDnn</dnn-redirection-rule>
</dnn-redirection-profile>
<ebm-data-options operation="replace">
  <include-gw-userplane-ip>true</include-gw-userplane-ip>
</ebm-data-options>
<geo-redundant-pool operation="replace">
  <ue-backup-distribution-option>wholePDUH</ue-backup-distribution-option>
  <periodic-backup-timer>123</periodic-backup-timer>
  <use-weighted-replication>false</use-weighted-replication>
  <skue-backup-if-cell-change-only>false</skue-backup-if-cell-change-only>
  <allow-second-retrieval>true</allow-second-retrieval>
</geo-redundant-pool>
<gtp-v2 operation="replace">
  <allow->
```





# Configuration Philosophy

Several unsolved problems introduce toil into our configuration management system.

- **Complexity in configurations**

5G core configurations typically involve thousands of parameters and are thus not user friendly. Abstraction reduces complexity by focusing on essential controls.

- **Redundancies across configurations**

Repeated parameters across application configuration and helm charts. Abstraction unifies these redundancies into a single source of truth.

E.g. ASN in a BGP setup needs to be known by both peers.

- **Limited configuration discovery**

Automatic configuration based on the environment still very weak.

E.g. Use of status field of a Kubernetes Resource as the input value of a Helm Release.



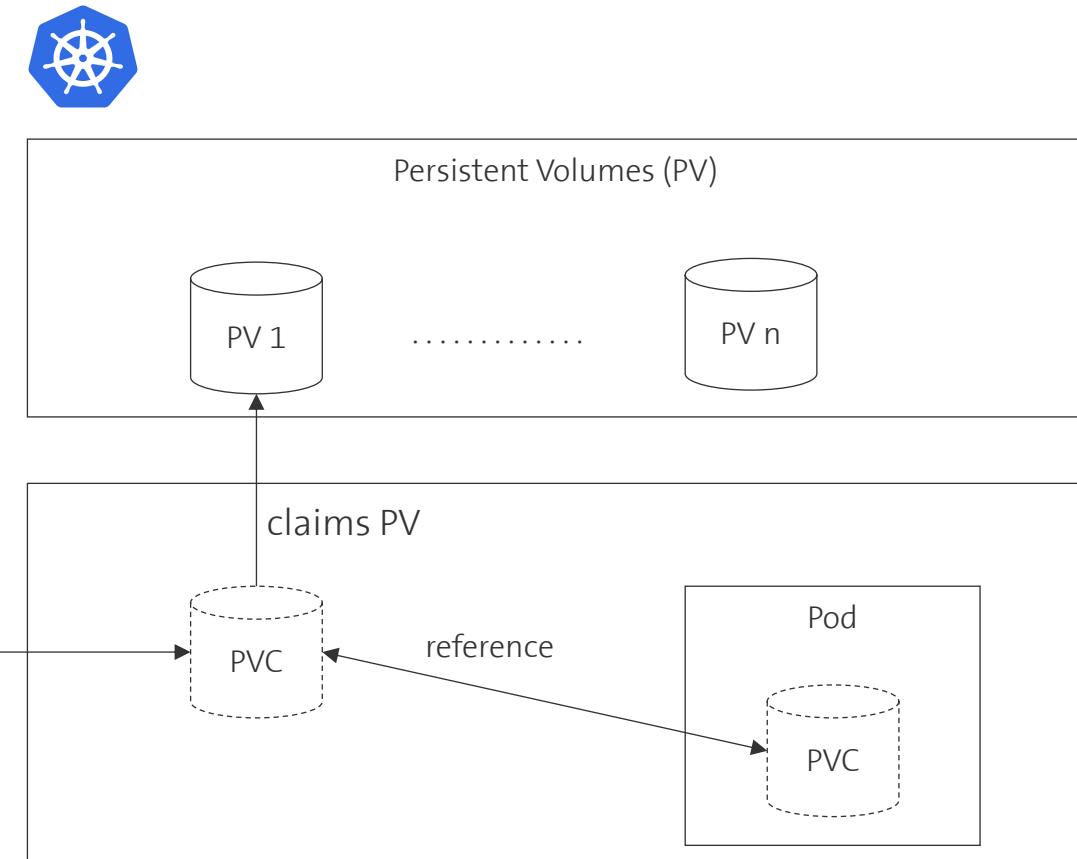
# Shared Source of Truth in Kubernetes and GitOps

Kubernetes claim model used in Persistent Volumes (PV) and Persistent Volume Claims (PVC)

- Decouples resource requests from provisioning
- PVCs represents user needs, PVs fulfill them
- Users interact using on claim names, not details like volume IDs
- Only the claims are stored in Git, details such as PV IDs are stored in Kubernetes
- Source of truth is split between Git (intent) and Kubernetes (dynamic state)
- Enables abstraction and reusability



Flux creates PVC defined in Git





# Benefits of Abstraction



## Reduce Toil

Simplifies tasks by focusing on essential controls, reducing manual interventions.



## Increased Reliability

Ensures consistent configurations across systems, minimizing errors and mismatches.



## Enhanced Efficiency

Streamlines configuration processes, saving time and resources.

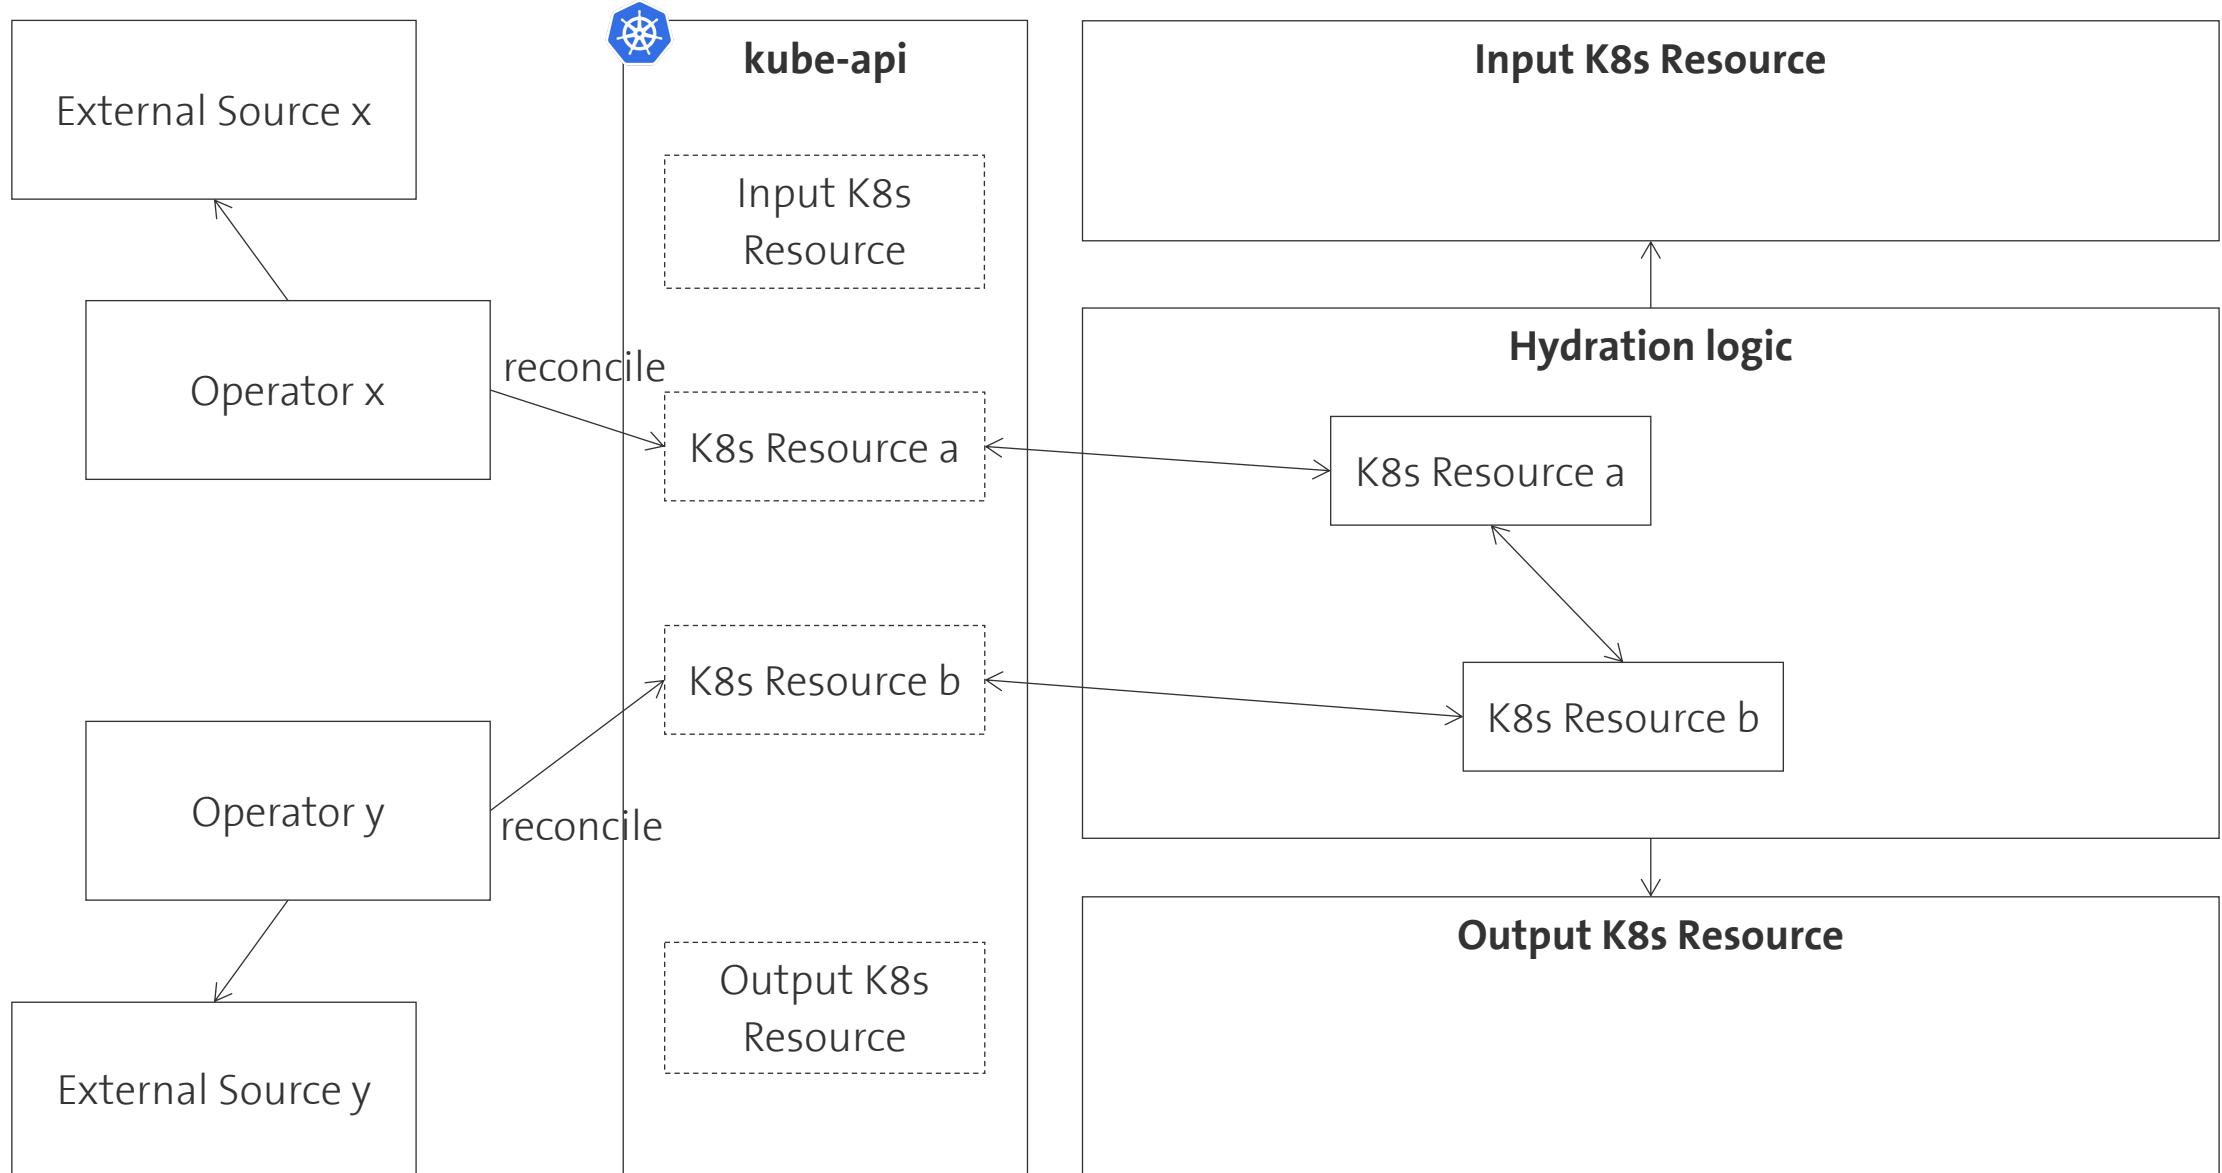


## Scalability

Facilitates easier scaling by managing configurations at a higher abstraction level.



# Flow of Config Abstraction





# What Are Our Requirements?



## KRM Based

Input and Output can be KRM resources



## Garbage Collection

Resources belonging together are lifecycled together



## Output Flexibility

Generated resources stored with various backends



## Reconciliation

Output is updated as Input is changed



## Functions

Ability to define helper functions (e.g. string manipulation)



## Gradual Adoption

Migration from static to dynamic taken at own pace.



## But What About Helm, Kustomize, Argo and Flux?



### Limited Dynamic Assembly

We cannot use live Kubernetes resources as source of information for e.g. a helm release



### No Custom Functions

We cannot invoke arbitrary code into Kustomize/Helm templating



### Only Partially KRM

Not all Resources involved follow KRM (e.g. Helm Values)



# Kubernetes Resource Model (KRM)



## API extensions

Custom Resource Definitions extend the Kubernetes API.

e.g. API definition for NetBox PrefixClaim



## CRs as Instances

Custom Resources instantiate a CRD.

e.g. NetBox PrefixClaim resource



## Business Logic

Use of Operators or templates to run custom logic

e.g. Assemble a config, Self healing

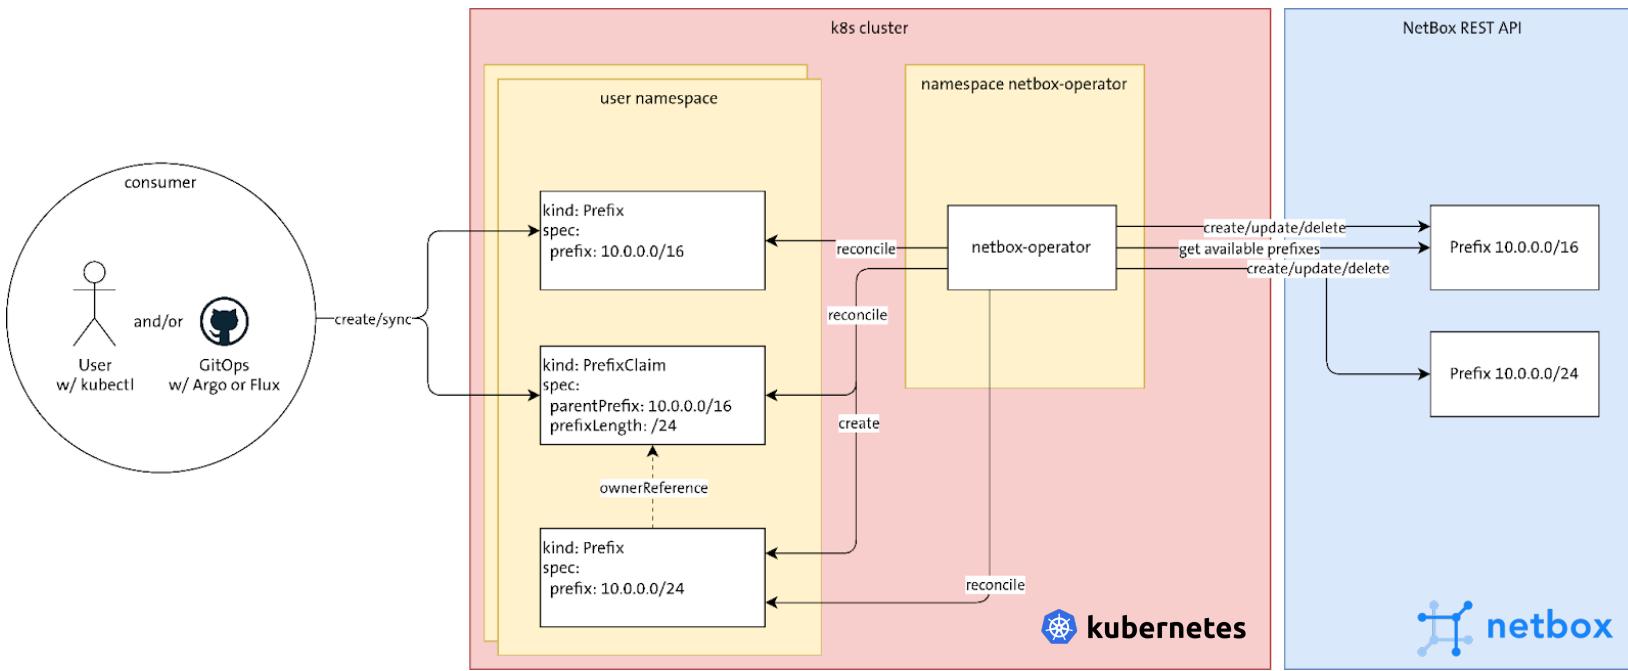


# NetBox Operator – Now Open Source

NetBox Operator, an open-source tool designed to integrate NetBox resource management directly into your Kubernetes environment.

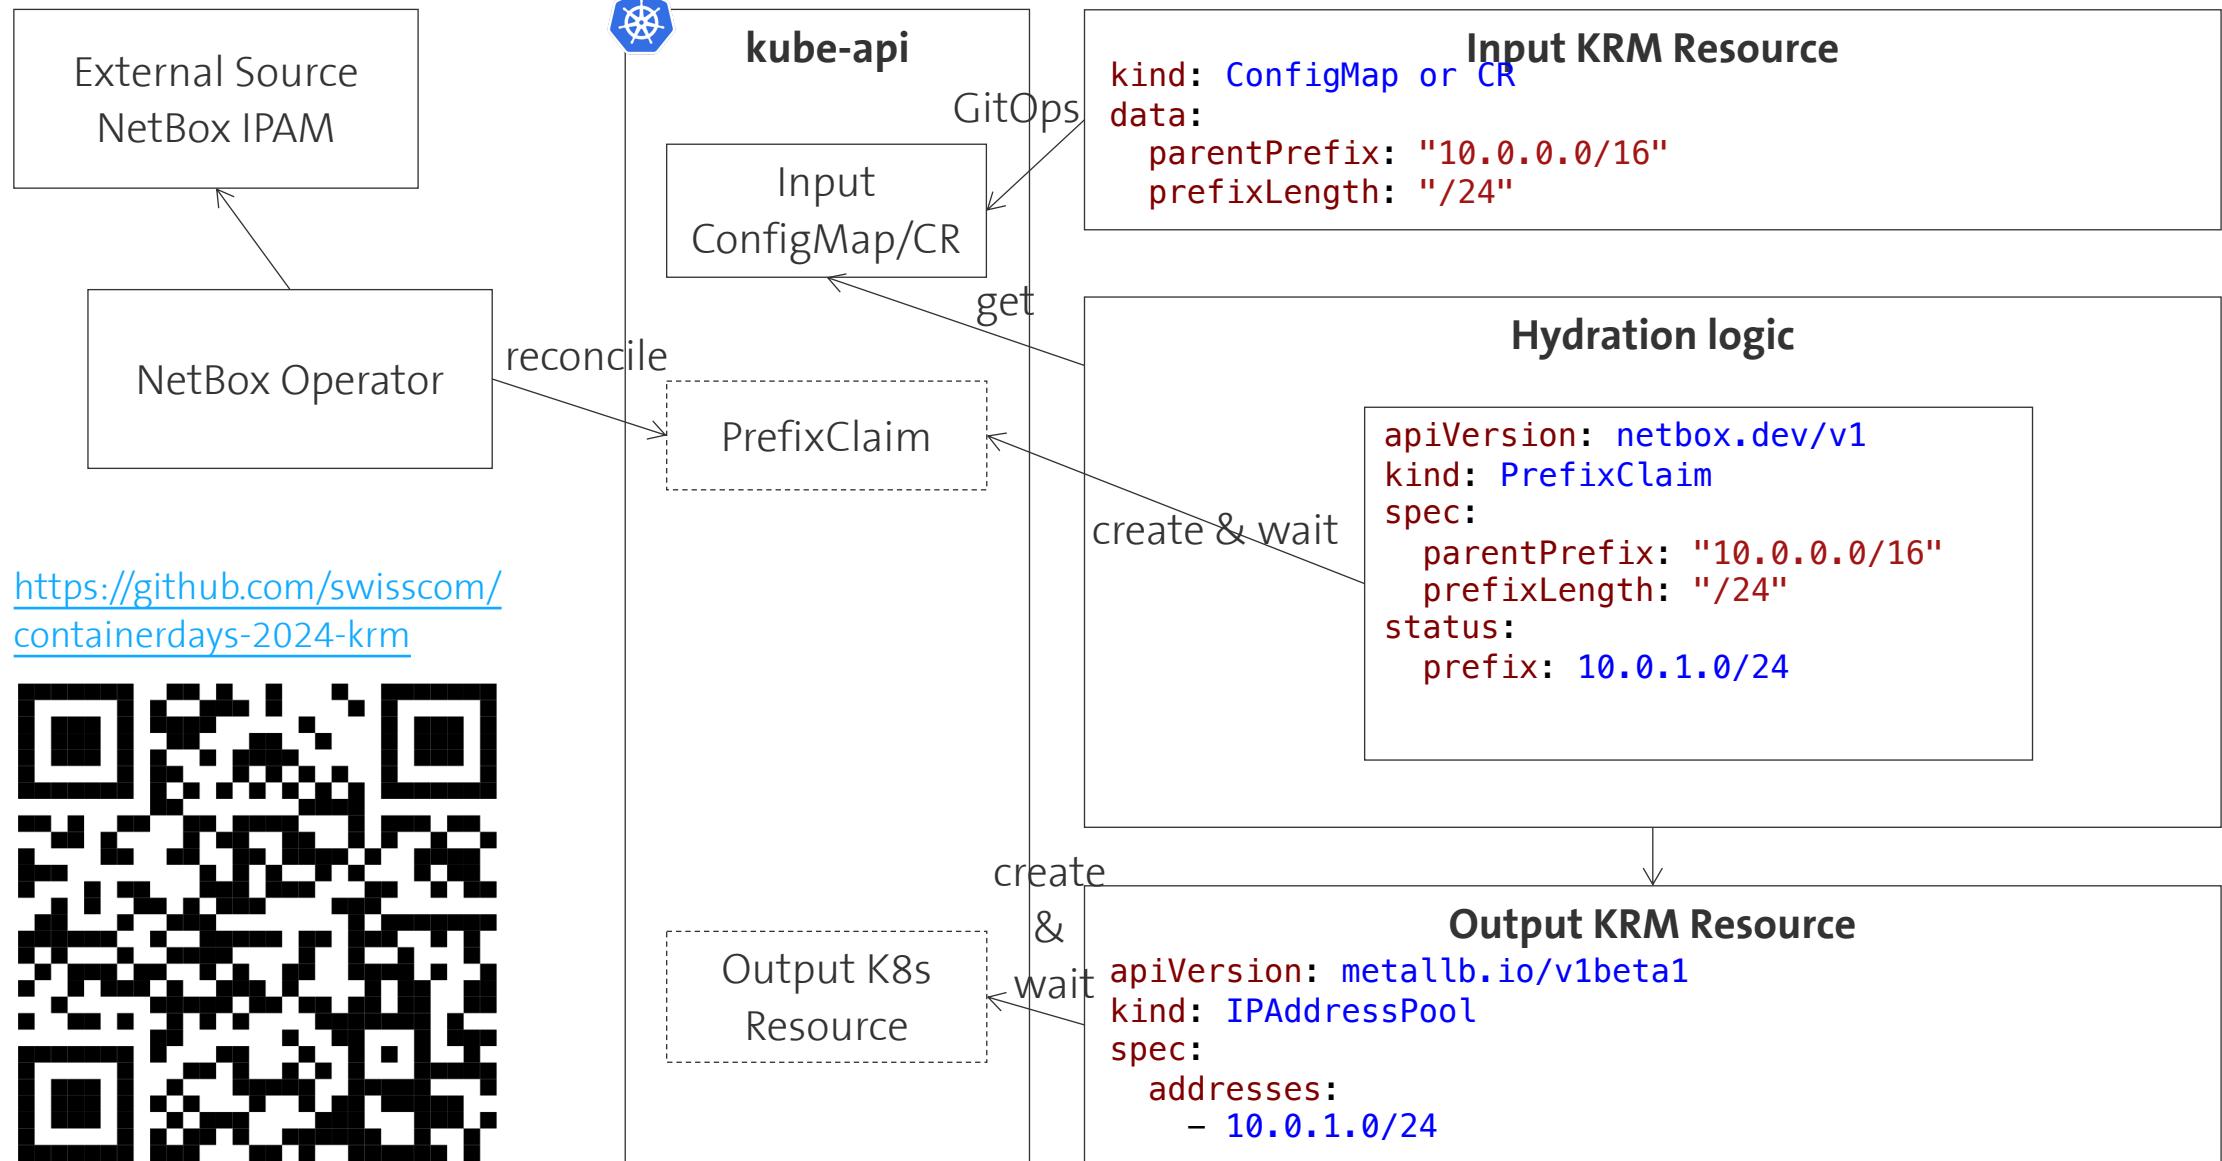
<https://github.com/netbox-community/netbox-operator>

[Announcement on www.swisscom.ch](#)





# Config Abstraction Demo Example: Dynamically Assign IPs for MetallB





# Ansible + Jinja2

**Open-Source automation framework**

<https://github.com/ansible/ansible> (62k Stars, 23k Forks)

**Business Logic in Templates and Playbooks**

Jinja2 for templating, Playbooks for imperative logic

**State Management**

Stateless, can use Kubernetes Garbage Collection

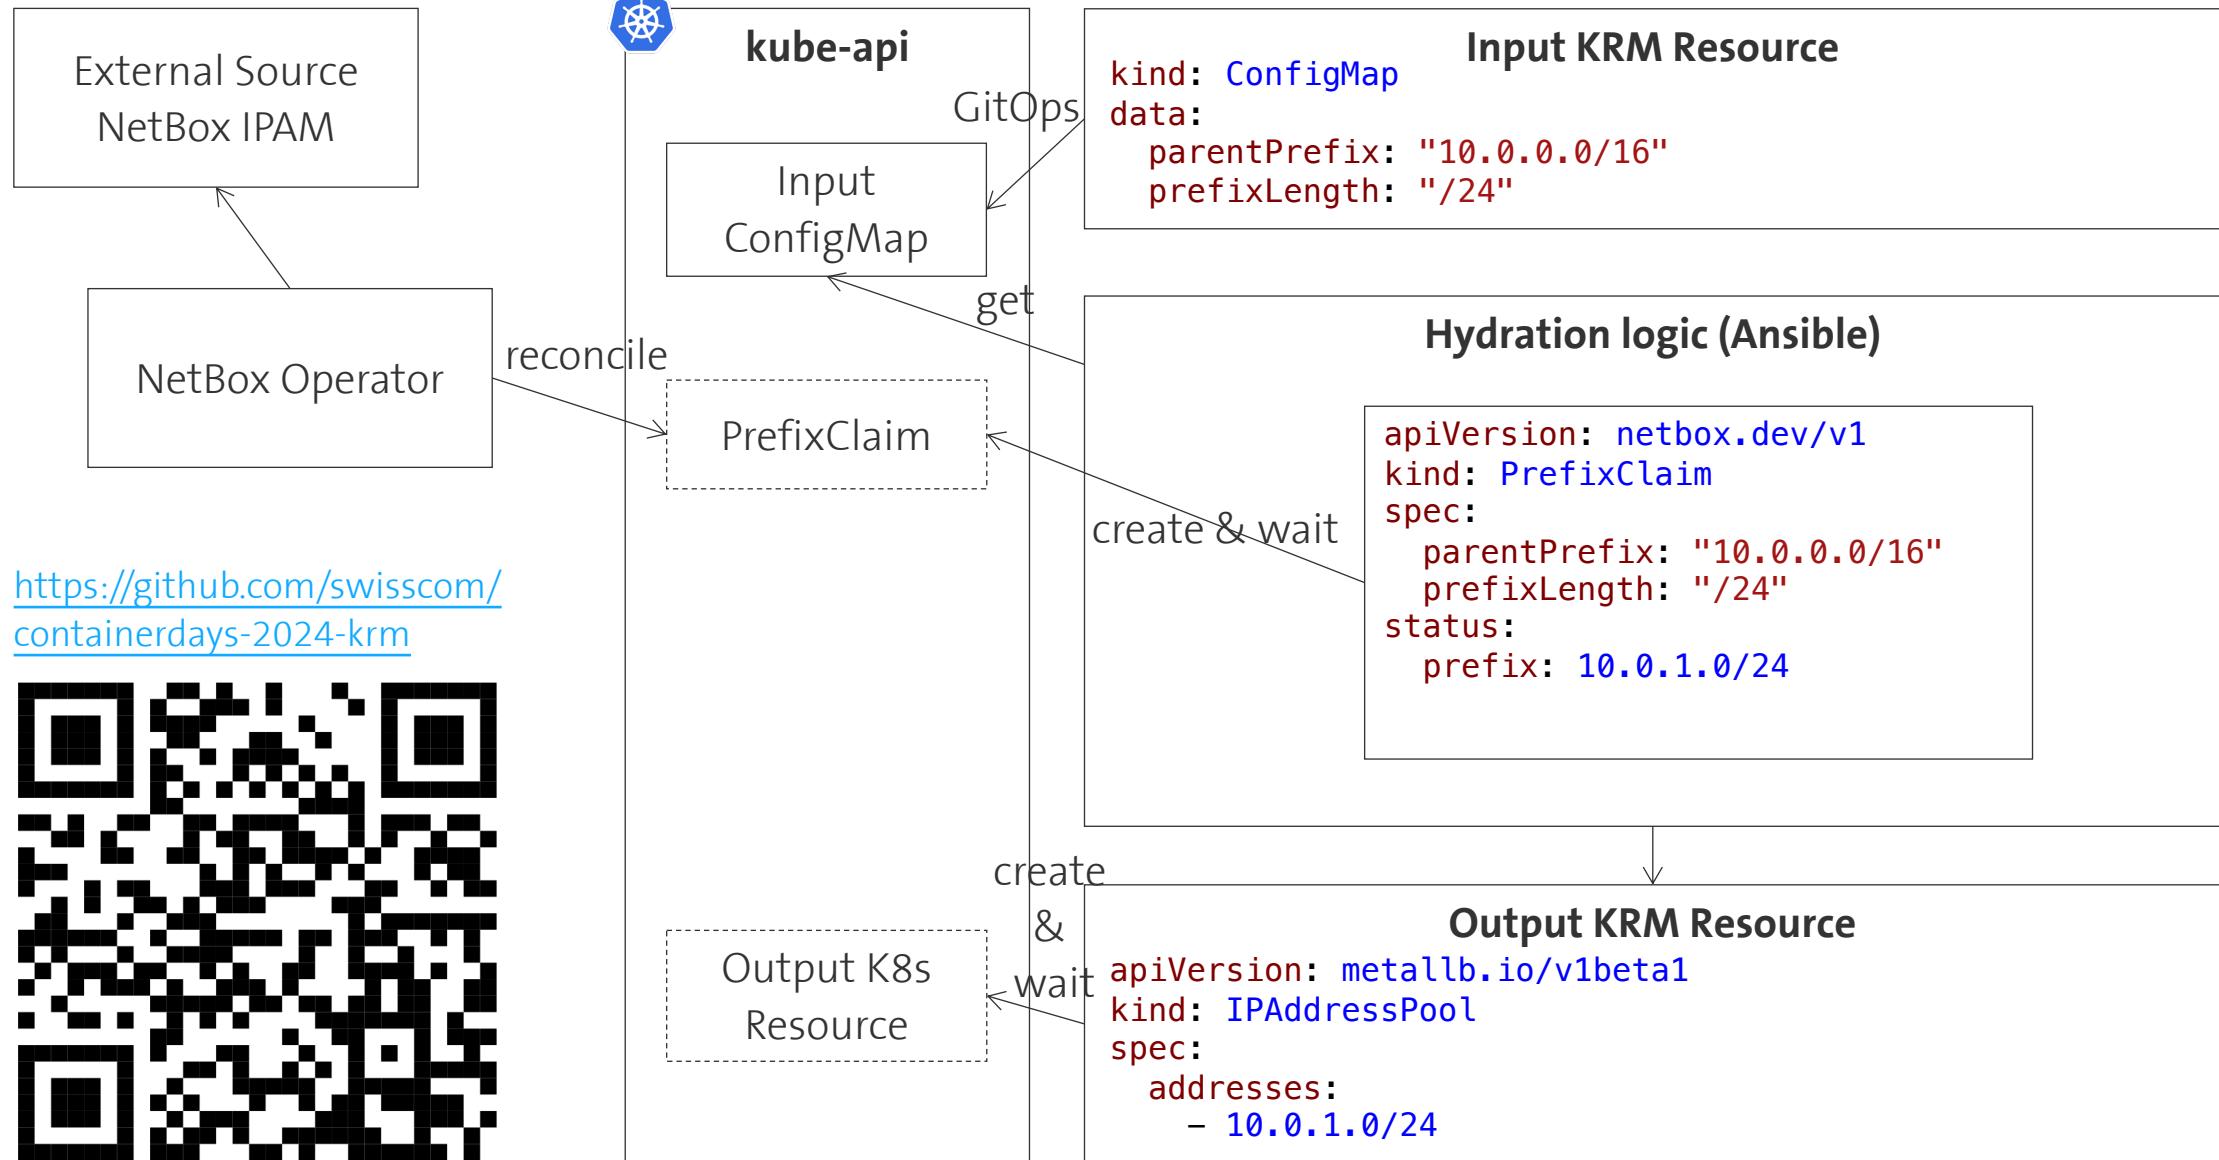
**No Reconciliation**

Just a binary, triggering/execution not solved





# Config Abstraction With Ansible + Jinja2





# KForm

**KRM templating inspired by Terraform**

<https://github.com/kform-dev/kform> (4 Stars, 1 Fork)

**Business Logic with minimum Code**

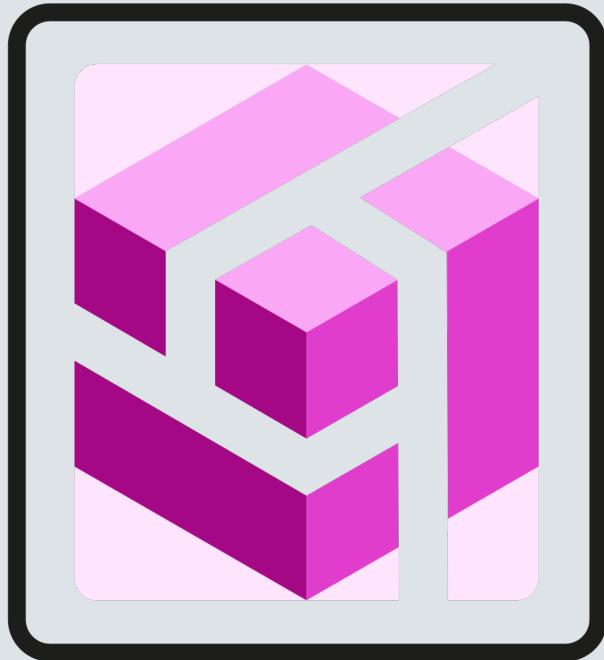
[cel-go](#) used for logic within Kubernetes Resources

**State Management**

Tracks state using ConfigMap in the kube-api

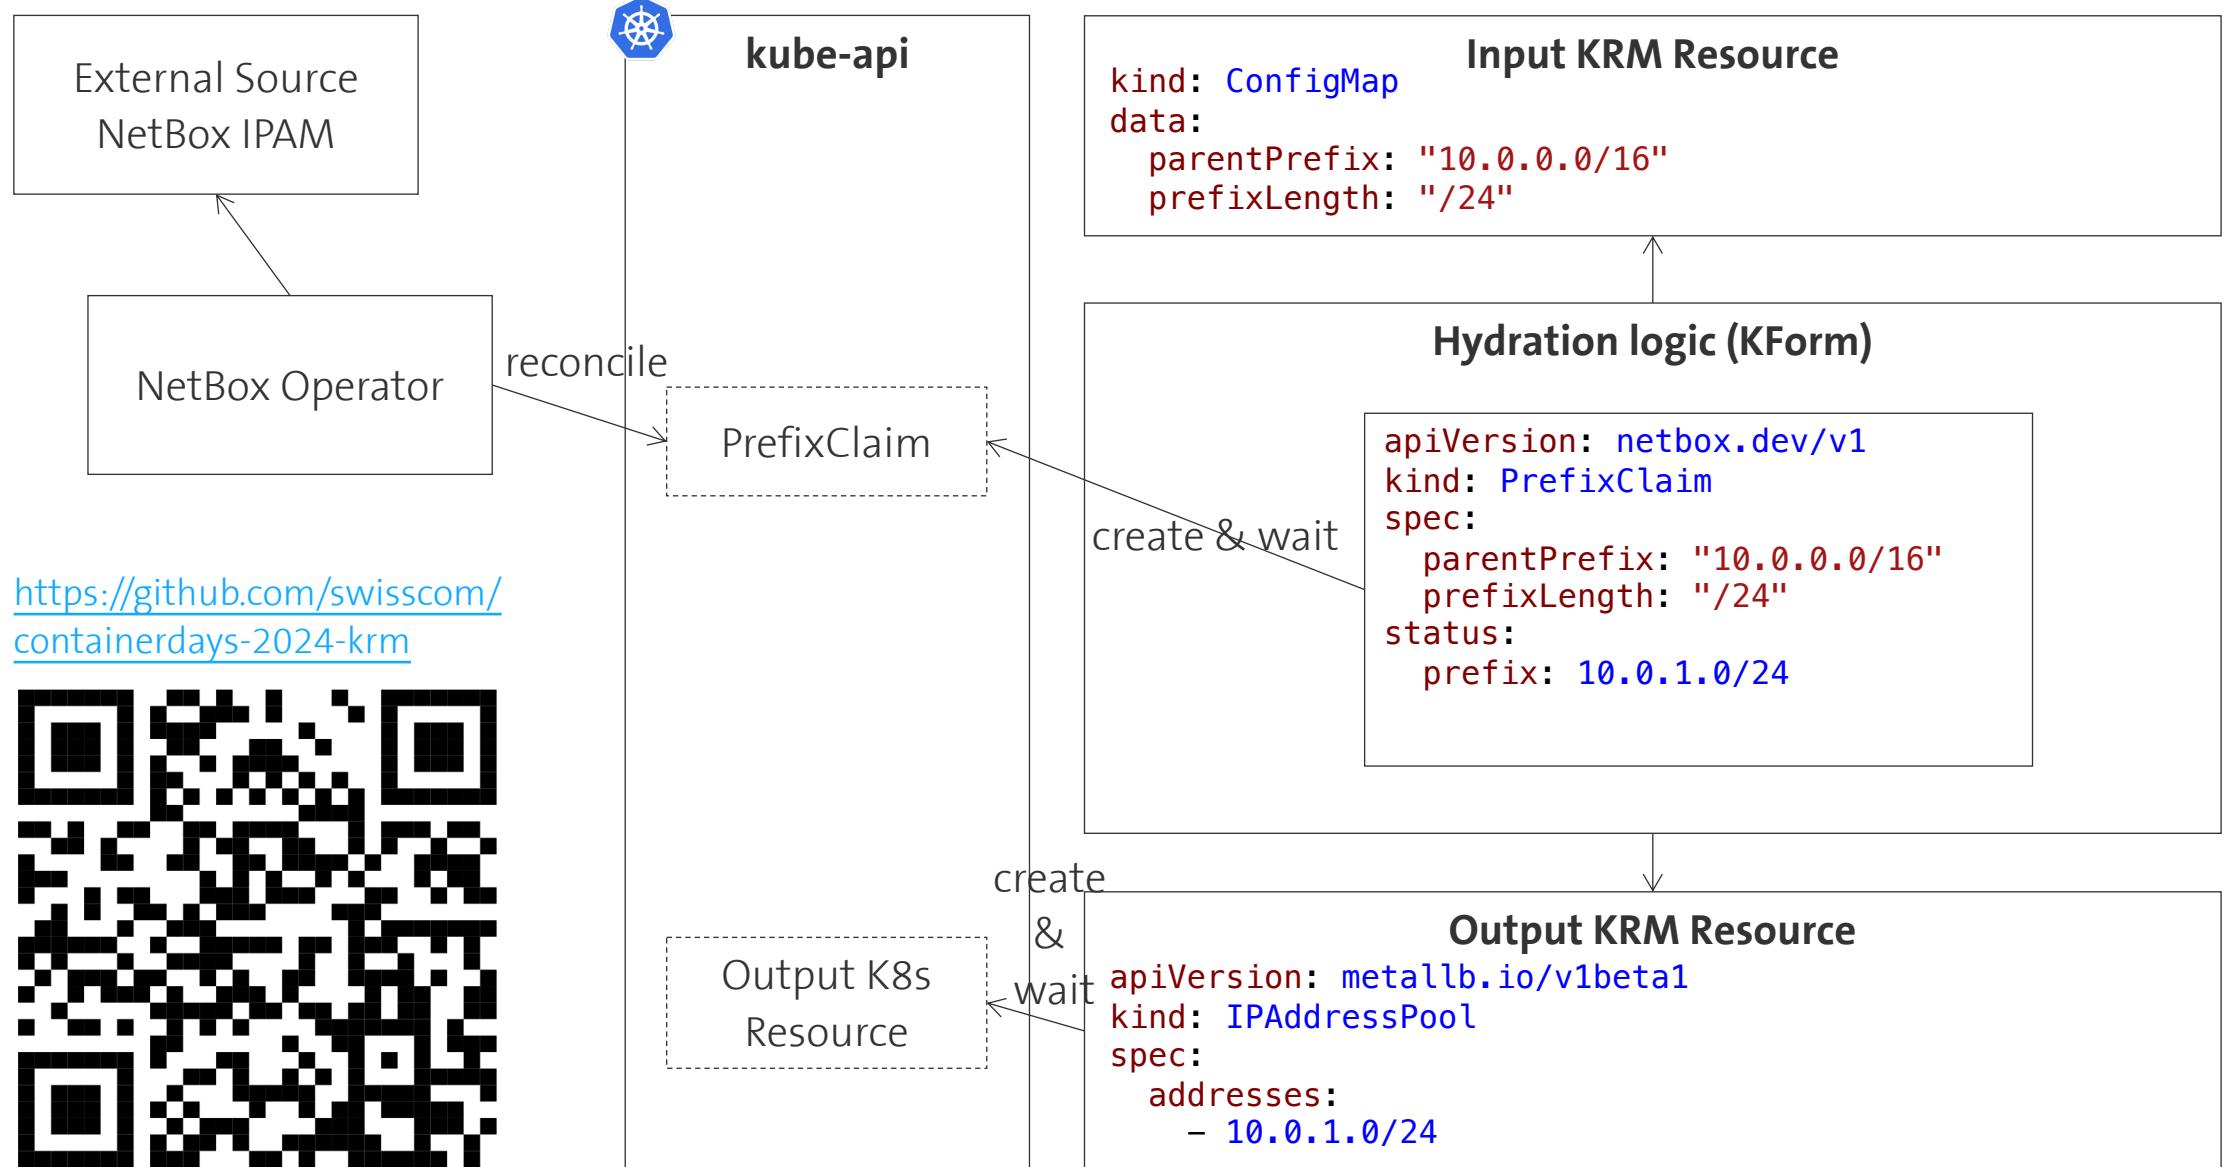
**No Reconciliation**

Just a binary, triggering/execution not solved





# Config Abstraction With KForm





# Kubernetes Operator

**Golang Based Operator scaffolded using Kubebuilder**

<https://github.com/kubernetes-sigs/kubebuilder>

**Business Logic in Go Code**

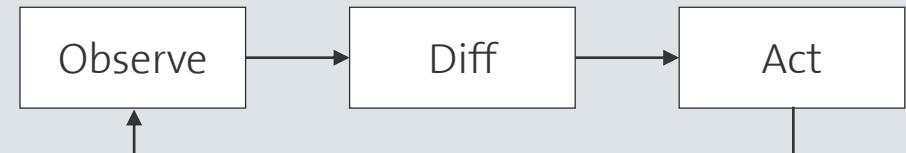
Maximum flexibility for imperative logic

**State Management**

Uses Kubernetes Garbage Collection and Finalizers

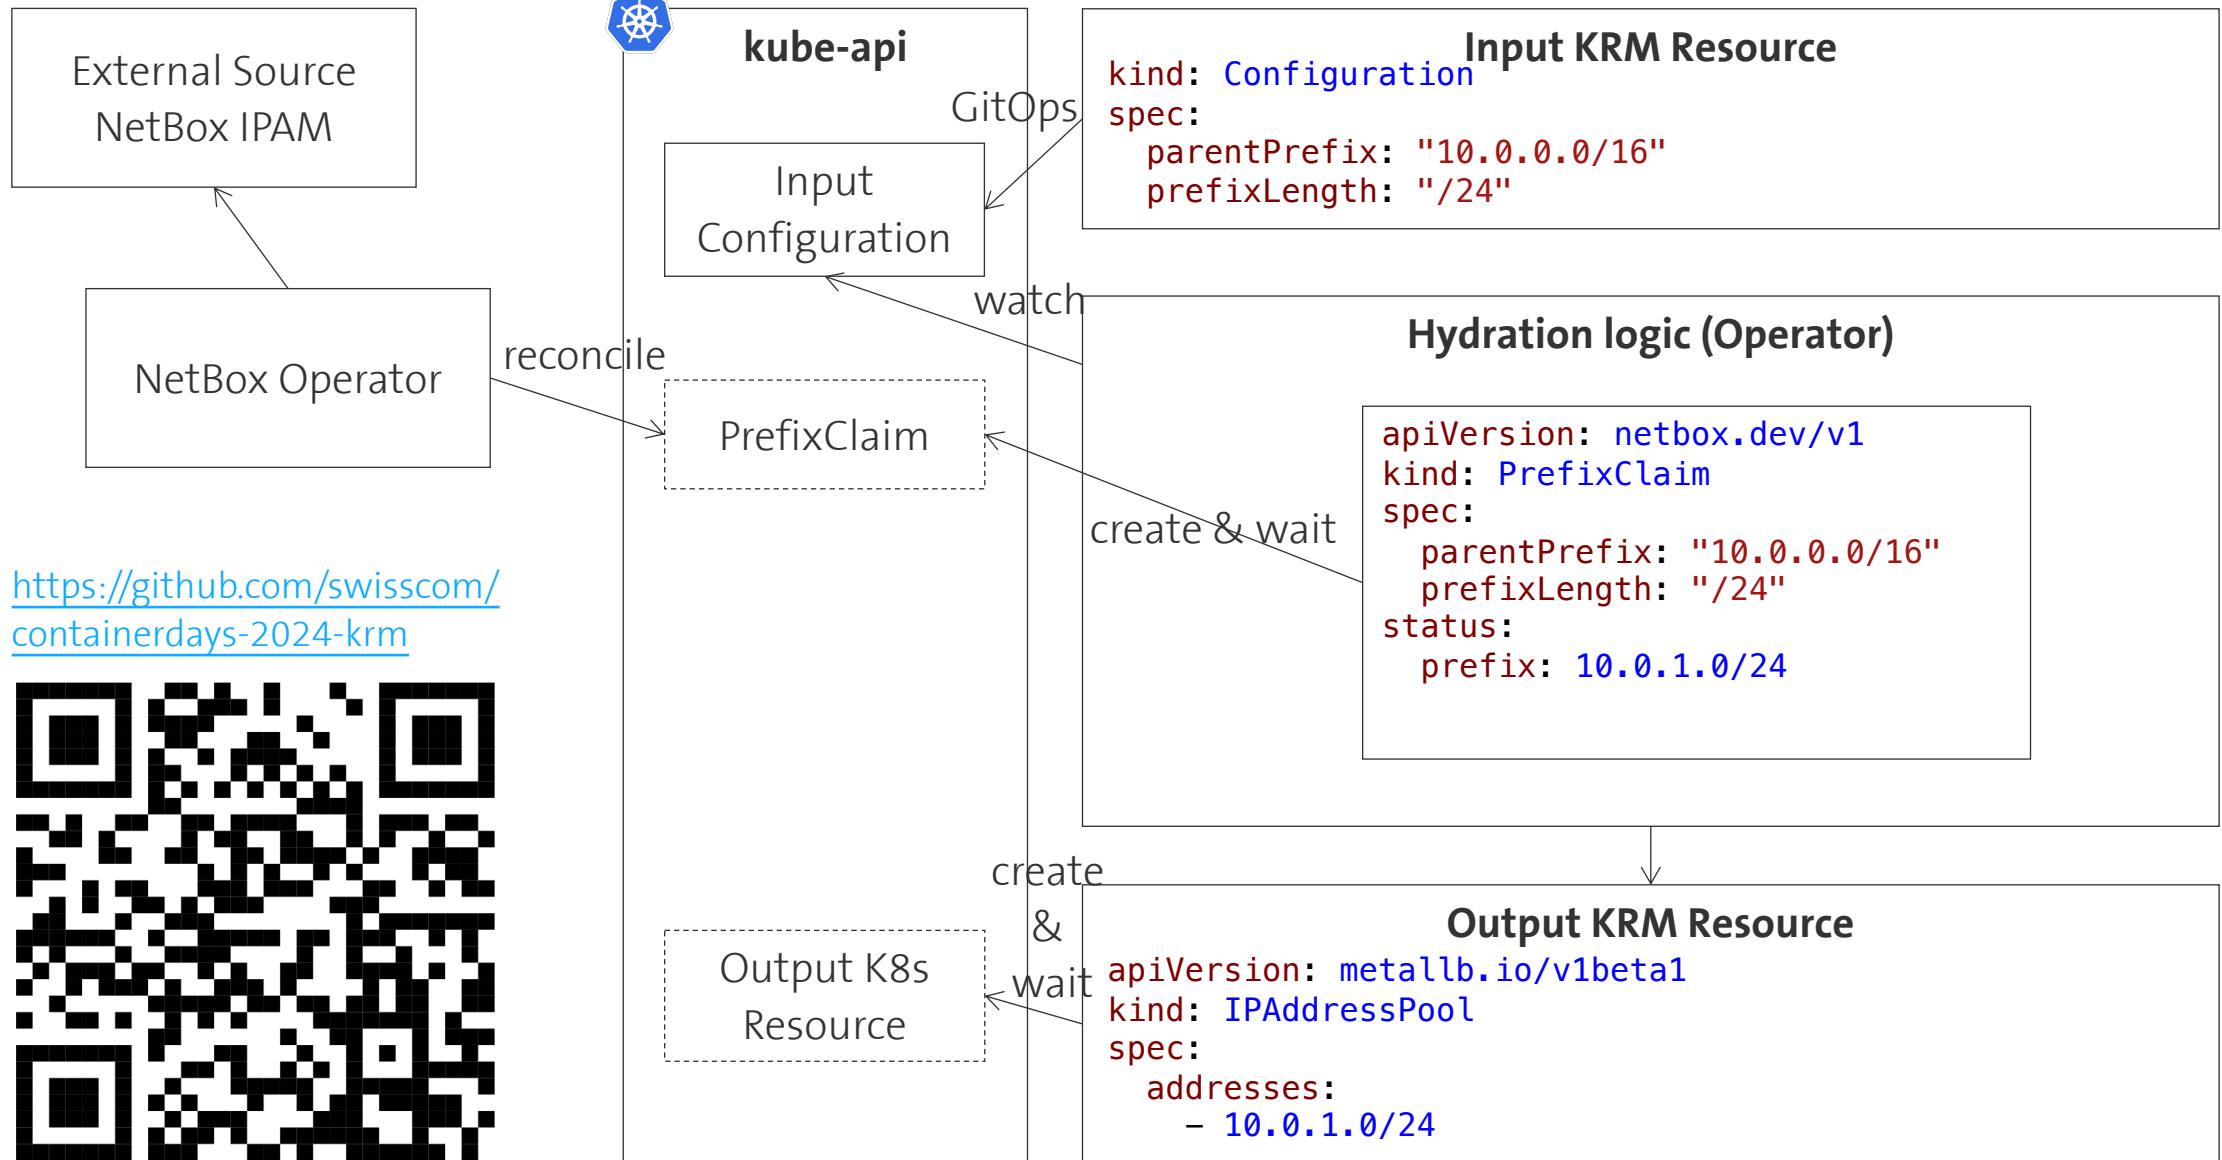
**Reconciliation**

Event driven triggering, constant reconciliation





# Config Abstraction With Kubernetes Operator





# Ansible vs. KForm vs. Kubernetes Operator

|                    | Ansible+Jinja2 | KForm     | Kubernetes Operator |
|--------------------|----------------|-----------|---------------------|
| KRM Based          | ● ● ○ ○ ○      | ● ● ● ○ ○ | ● ● ○ ○ ○           |
| Garbage Collection | ● ● ● ● ○      | ● ● ● ○ ○ | ● ● ● ● ○           |
| Output Flexibility | ● ● ● ● ●      | ● ● ○ ○ ○ | ● ● ● ● ●           |
| Reconciliation     | ● ○ ○ ○ ○      | ● ○ ○ ○ ○ | ● ○ ○ ○ ○           |
| Functions          | ● ● ● ● ○      | ● ○ ○ ○ ○ | ● ● ● ● ○           |
| Gradual Adoption   | ● ● ● ● ●      | ● ● ● ○ ○ | ● ○ ○ ○ ○           |

**No mature Kubernetes native tool available to dynamically hydrate resources**



# Key Takeaways



## Dynamic Parameter Management

Difficult to dynamically add parameters to Kubernetes resources in GitOps



## Tooling Limitations

No mature tools for config hydration with KRM



## Existing Solutions' Shortcomings

**Ansible & Jinja:** Powerful templating, but not designed with K8s in mind

**KForm:** KRM-based but lacks scalability for multiple instances

**Golang Operator:** Too specific, leading to operational overhead



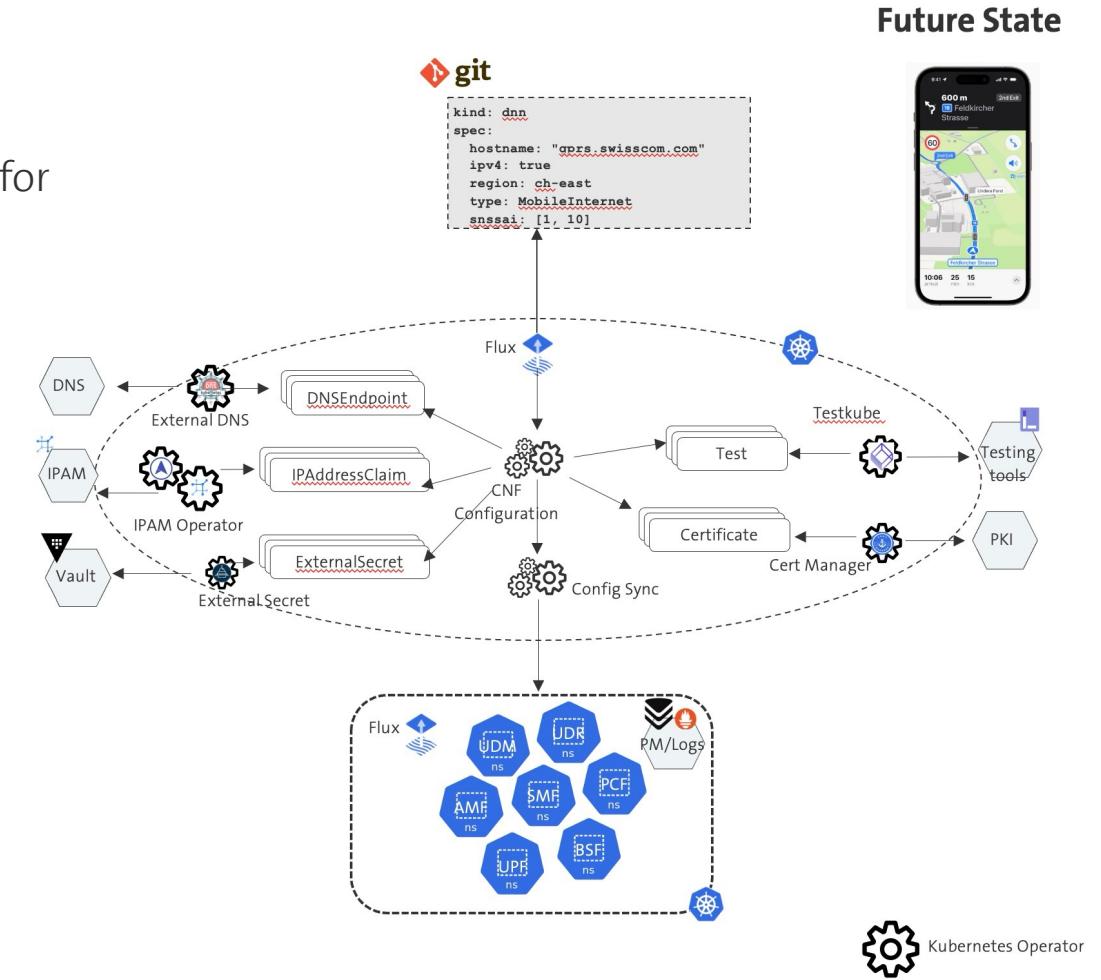
## Industry-Wide Relevance:

These challenges are not exclusive to the Telco industry



# Next Steps

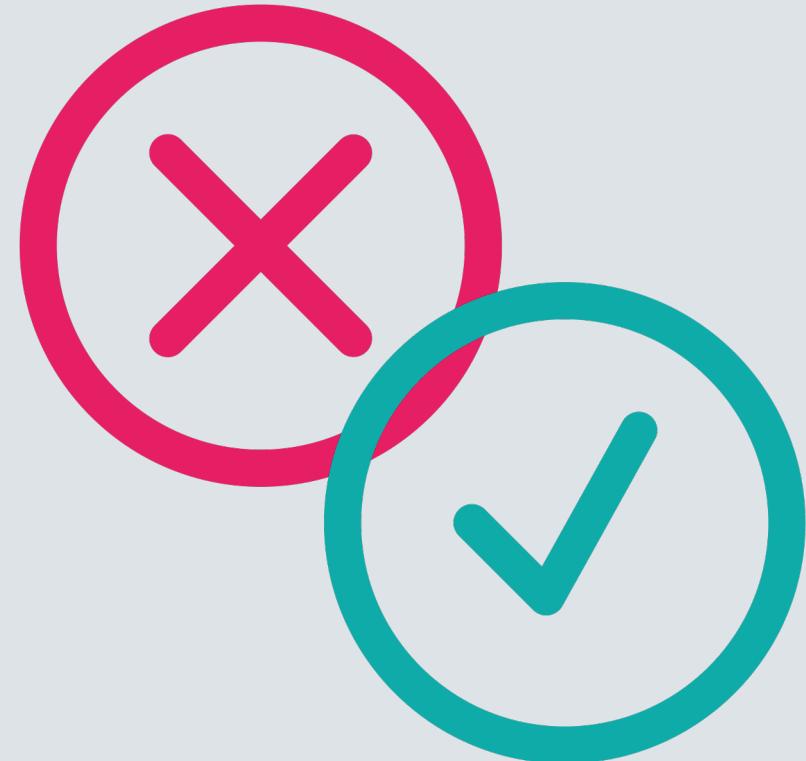
- **Monitor and Explore:** Track kpt, Kform, and other advancements in Kubernetes management
- **Operator Development:** Implement Kubernetes operators for generic templating
- **Multi-cluster management:** Enhance multi-cluster management and explore cloud-cluster tools (e.g. Nephio, KubeNet)
- **Safe Rollouts:** Ensure safe rollouts with configuration versioning
- **Collaboration:** Engage with open-source community and vendors





# Our Learnings & Suggestions

... on configuration abstraction





# Avoid

## **Checking in low level configurations in Git**

Things like IP Addresses, VLAN IDs

## **Complicated configurations**

Avoid unnecessary layers of abstraction

## **Ignore tooling gaps**

Avoid using tools like Ansible or Jinja that aren't designed with Kubernetes in mind

## **Neglect Industry Relevance**

Don't assume KRM challenges and solutions are unique to your sector





## Aim to

### Leverage abstraction

Simplify complex configurations by focusing on essential controls

### Reuse cloud native tools to be in-band with K8s

E.g. Flux, Argo, Testkube, cert-manager

### Prioritize scalability

Design KRM-based tools with future growth and scaling

### Stay Agile

Continuously refine and adapt your KRM practices to evolving needs and technologies

### Contribute to the ecosystem

Share your code





## Related Talks

### **Building and Operating a Highly Reliable Cloud Native DNS Service With Open Source Technologies**

*by Fabian Schulz, Hoang Anh Mai and Joel Studler at ContainerDays 2024*

Stage K6, Today 17:30 - 18:05

### **How We Are Moving from GitOps to Kubernetes Resource Model in 5G Core**

*by Ashan Senevirathne and Joel Studler at KubeCon Europe 2024*

<https://www.youtube.com/watch?v=crmTnB6Zwt8>



### **From Spreadsheets to "Everything as Code"**

*by Joel Studler and Jérémie Weber at ContainerDays 2023*

<https://www.youtube.com/watch?v=eMAMqW2aLTI>





# Thanks!





# Q&A





# GitOps vs. GitOps with KRM

|                                 | GitOps                                     | GitOps with KRM                                           |
|---------------------------------|--------------------------------------------|-----------------------------------------------------------|
| <b>Source of Truth</b>          | Git is the single source of truth          | Split between Git (intent) and Kubernetes (dynamic state) |
| <b>Configuration Management</b> | All details are in Git                     | Git holds high-level configs; Kubernetes manages the rest |
| <b>Abstraction</b>              | Users manage all details                   | Users focus on intent; Kubernetes abstracts complexity    |
| <b>Flexibility</b>              | Less flexible, changes require Git updates | More flexible, Kubernetes adjusts resources dynamically   |
| <b>Scalability</b>              | Can be complex to scale                    | Easier to scale with Kubernetes managing dynamics         |
| <b>Operational Complexity</b>   | High; all details tracked in Git           | Lower; Kubernetes automates dynamic management            |